
On Improving Decentralized Hysteretic Deep Reinforcement Learning

Xueguang Lu¹ Christopher Amato¹

Abstract

Recent successes of value-based multi-agent deep reinforcement learning employ optimism in value function by carefully controlling learning rate (Omidshafiei et al., 2017) or reducing update probability (Palmer et al., 2018). We introduce a decentralized quantile estimator: Responsible Implicit Quantile Network (RIQN), while robust to teammate-environment interactions, able to reduce the amount of imposed optimism. Upon benchmarking against related Hysteretic-DQN (HDQN) and Lenient-DQN (LDQN), we find RIQN agents more stable, sample efficient and more likely to converge to optimal policy.

1. Introduction

Recent development in multi-agent reinforcement learning (MARL) have borrowed increasing amount insights from single-agent RL developments, including deep convolutional networks for solving tasks high dimensional sensory observations (Mnih et al., 2015). In this paper, we bring insight from recent development in Deep Distributional RL to multi-agent settings; aim to be robust to sub-optimal local policies, we utilize estimated distributions to adjust learning rates in a more effective manner. More specifically, we apply Implicit Quantile Network [3] to multi-agent settings, along with a modified version we call Responsible-IQN.

Our work focuses on fully cooperative Independent Learners which both learn and execute in a distributed manner. This complete decentralization setting poses a harder problem than settings having information, typically actions, shared across agents (Claus & Boutilier, 1998). Ideally, the agents should able to robustly converge to joint globally optimal policy instead of locally optimal policy under local actions and local observations which subjects to environment non-stationary caused by actions taken by other agents (Fulda & Ventura, 2007). In other words, as agents are expected

to perform cooperative tasks, agent must be robust to often occurring non-effective explorations. Recent attempts on solving the issue in high dimensional observation settings have shown success with hysteretic Q-Learning (Matignon et al., 2007) and leniency (Panait et al., 2006). Both approaches limit negative value updates, either proportionally or probabilistically, aiming to partly ignore the effect of noneffective explorations.

Empirically, leniency have higher performance compared to hysteretic methods due to its adaptive leniency values at different stages of training (Palmer et al., 2018). Our approach, based on hysteretic Q-Learning, adapts to different training stage in state-action space, resulted in more stable convergence towards optimal joint policy.

2. Background

2.1. Markov Decision Process and Deep Q-Network

We formulate perfect information sequential decision making problem as a Markov Decision Process $\langle \mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R} \rangle$ where \mathcal{S} is a finite state space, \mathcal{A} a finite action space, $\mathcal{T}(s, a, s')$ the probability of transitioning from state $s \in \mathcal{S}$ to $s' \in \mathcal{S}$ by taking action $a \in \mathcal{A}$, and $\mathcal{R}(s, a, s')$ is the immediate reward for such transition. The Markovian assumption is implied by the transition probability. Under this formulation, the problem is to find a optimal policy $\pi^* : \mathcal{S} \rightarrow \mathcal{A}$ which maximize sum of rewards.

Q-Learning (Watkins & Dayan, 1992) is a popular model-free method of learning optimal policy, which iterate on a set of Q values or approximators to approach optimal Q values or estimates, where $Q(s, a)$ is the expected maximum sum of rewards achievable in the future given in state s and taking action a .

$$Q^*(s, a) = \max_{\pi} \mathbb{E}(r_t + \gamma r_{t+1} + \dots | s_t = s, a_t = a, \pi) \quad (1)$$

When used with approximation, tabular Q-Learning methods often guarantees to convergence on discrete set of states and actions. The update of tabular approach is shown below.

$$Q(s_t, a_t) \leftarrow (1 - \alpha)Q(s_t, a_t) + \alpha(r_t + \gamma \max_a Q(s_{t+1}, a)) \quad (2)$$

Where $\alpha \in (0, 1]$ is learning rate at which the Q value

¹College of Computer and Information Science, Northeastern University, Boston, MA. Correspondence to: Xueguang Lu <lu.xue@husky.neu.edu>, Christopher Amato <c.amato@northeastern.edu>.

updates, and $\gamma \in (0, 1]$ is the discount factor applied to future Q estimations.

Deep D-Network (Mnih et al., 2015) develops on a common practice where a function approximator is used for estimating Q values by parameterize Q function $Q(s, a, \theta)$ with parameters θ using deep convolutional neural network. DQN uses experience replay (Lin, 1993) where each transition (s_t, a_t, r_t, s_{t+1}) is stored into a large fix-sized experience buffer $D_t = \{(s_1, a_1, r_1, s_2), \dots, (s_t, a_t, r_t, s_{t+1})\}$ from which all training batches for the network are uniformly sampled. The update of the network follows the following loss function:

$$L_i(\theta_i) = \mathbb{E}_{s,a,r,s' \sim U(D)} [(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i))^2] \quad (3)$$

where θ_i^- is the parameters for target network at iteration i , a target network is an identical network whose parameters are never directly updated but copied from the main network every C steps in order to maintain value stability.

2.2. Partial Observability

Single-agent tasks without access to perfect information often formalized as Partial Observable MDP (POMDP) problems. POMDP extends MDP by granting agent observations instead of states: $\langle \mathcal{S}, \mathcal{A}, \mathcal{Z}, \mathcal{T}, \mathcal{O}, \mathcal{R} \rangle$; where \mathcal{S} is still a finite state space, \mathcal{A} is finite action space, \mathcal{Z} is finite set of observations, $\mathcal{T}(s, a, s')$ is the probability of transitioning from state $s \in \mathcal{S}$ to $s' \in \mathcal{S}$ by taking action $a \in \mathcal{A}$, $\mathcal{O}(s, a, z)$ is the probability of observing $z \in \mathcal{Z}$ given state s and action a , $\mathcal{R}(s, a)$ is the immediate reward for such transition.

Extending DQN to accommodate partially observable tasks, (Hausknecht & Stone, 2015) proposed a model called Deep Recurrent Q-Network (DRQN), where a recurrent, LSTM (Hochreiter & Schmidhuber, 1997) in particular, layer was used to replace the first post-convolutional fully-connected layer of DQN. Hausknecht and Stone argues that the recurrent layer is able to integrate an arbitrarily long history which can be used to infer the underlying state. Empirically, DRQN out-performed DQN on partial-observable tasks and is on par with DQN on fully-observable tasks. Formally, the network estimates $Q(o_t, a_t, h_{t-1} | \theta)$ (instead of $Q(s_t, a_t, \theta)$), where θ is the parameters of the network, and $h_{t-1} = LSTM(h_{t-2}, o_t)$ is the hidden state generated from the LSTM layer from $t - 1$ time step. Although the model arguably appears to conform more closely neuroscientifically, Hausknecht and Stone nonetheless points out that DRQN confers no systematic benefit compared to DQN where past k observations is used to approximate s .

2.3. Decentralized POMDP (Dec-POMDP)

Dec-POMDP aims to generalize POMDP to multi-agent settings, as actions and observations becomes joint actions

and observations: $\langle \mathcal{I}, \mathcal{S}, \mathbf{A}^{\mathcal{I}}, \mathcal{Z}, \mathcal{T}, \mathbf{O}^{\mathcal{I}}, \mathcal{R} \rangle$ where \mathcal{I} is a finite set of agents, \mathbf{A}^i is action space for agent $i \in \mathcal{I}$, and \mathbf{O}^i is observation space of agent i . At every time step, a joint action $\mathbf{a} = \langle a^1, \dots, a^{|\mathcal{I}|} \rangle$ is taken, and agents receive a shared immediate reward $\mathcal{R}(s, \mathbf{a})$.

3. Related Work

3.1. Distributional RL and Implicit Quantile Network (IQN)

As a distributional RL method, quantile networks are interested in distribution over returns Z^π , where $\mathbb{E}(Z^\pi) = Q^\pi$, by estimating the inverse c.d.f. of Z^π , denoted F_π^{-1} . Implicit Quantile Network (IQN)[3] reparameterizes samples from a base distribution, e.g. $\tau \sim U([0, 1])$, where τ represents quantile value:

$$Q_\beta(x, a) := \mathbb{E}_{\tau \sim U([0, 1])} [F_{\pi, \beta(\tau)}^{-1}(x, a)] \quad (4)$$

[TODO: QUANTILE ENBEDDINGS] where $\beta : [0, 1] \rightarrow [0, 1]$ distorts risk sensitivity, risk neutrality is achieved when $\beta = \text{id}$.

The quantile regression loss[4] for quantile at τ and error e is defined as

$$\rho_\tau(e) = (\tau - \mathbb{I}\{e \leq 0\})e \quad (5)$$

which weighting overestimation by $1 - \tau$ and underestimation by τ .

Given two samples $\tau, \tau' \sim U([0, 1])$ and policy π_β , the sampled TD error for time step t is

$$e_t^{\tau, \tau'} = r_t + \gamma F_{\tau'}^{-1}(x_{t+1}, \pi_\beta(x_{t+1})) - F_\tau^{-1}(x_t, a_t) \quad (6)$$

Thus, given $\tau_{1:N}, \tau'_{1:N'}$, the loss becomes:

$$L = \frac{1}{N'} \sum_{i=1}^N \sum_{j=1}^{N'} \rho_{\tau_i}(e^{\tau_i, \tau'_j}) \quad (7)$$

[TODO: HOW DISTRIBUTION IS ESTIMATED THROUGH MULTIPLE RUNS]

3.2. Hysteretic Q learning

Due to non-stationary nature of multi-agent tasks caused by unobservable teammate policy, which is often sub-optimal due to exploration strategy, vanilla Q-Learning therefore would learn to estimate said non-stationary. HDQN[1] uses two learning rates $0 < \beta < \alpha \leq 1$, α is used for overestimation TD updates, and β is used otherwise. Overestimation is defined as $\delta_t > 0$; where TD error $\delta_t = r + \gamma \max_{a'} Q(s_{t+1}, a) - Q(s_t, a_t)$.

4. Methods

We propose a more granular approach for controlling the learning rate based on time difference quantile similarity measures. The measure, which we call TD likelihood, measures the overlapping area size of two probability density functions approximated by quantiles. The aim is that, for overall similar functions, even with drastic difference in specific quantile locations, the learning rate remains high, thus removing a portion of optimism imposed by hysteretic and lenient methods. We develop on IQN architecture, and optionally add LSTM layer for partial observable tasks.

4.1. Time Difference likelihood

We first introduce an approximation method for estimating the likelihood of samples $F_{\tau_{1:M'}}^{-1}(x_t, a_t)$, denoted $t_{1:M'}$, given a distribution constituted by set of samples $F_{\tau_{1:M}}^{-1}(x_t, a_t)$, denoted $d_{1:M}$. Let the distribution that $d_{1:M}$ approximates be $\mathcal{F}(X) = F_i(X)$ if $d_i \leq X \leq d_{i+1}$ otherwise 0, where each F_i is a linearity that fits (τ_i, d_i) and (τ_{i+1}, d_{i+1}) . We denote $\mathcal{P}(X) = P(X | d_{1:M})$ for notation simplicity. We observe that, for arbitrary a and b

$$\begin{aligned} \mathcal{P}(a < X < b) &= \int_a^b Z_{d_{1:M}}(X) dX \\ &= \int_{-\infty}^b Z_{d_{1:M}}(X) dX - \int_{-\infty}^a Z_{d_{1:M}}(X) dX \\ &= \mathcal{F}(b) - \mathcal{F}(a) \\ &= \sum_{i=1}^{M-1} \frac{|(a, b] \cap (d_i, d_{i+1}]|}{d_{i+1} - d_i} (\mathcal{F}(d_{i+1}) - \mathcal{F}(d_i)) \\ &= \sum_{i=1}^{M-1} \frac{|(a, b] \cap (d_i, d_{i+1}]|}{d_{i+1} - d_i} (\tau_{i+1} - \tau_i) \end{aligned}$$

Then, we can estimate target set likelihood as:

$$l_{t_{1:M'}, d_{1:M}} = \mathbb{E}_{j \in 1:M'} \mathcal{P}(\mathbb{E}(t_{j-1}, t_j) \leq t_j \leq \mathbb{E}(t_{j+1}, t_j)) \quad (8)$$

The likelihood measure differs from KL divergence and wasserstein distance

4.2. Responsible Hysteretic update

We introduce Responsible Hysteretic IQN by incorporating sample likelihood discussed above into hysteretic learning rate tuning. While traditional hysteretic uses learning rates $0 < \beta < \alpha \leq 1$, we use $0 < \max(\beta, l_{t_{1:M'}, d_{1:M}}) \leq \alpha \leq 1$.

More specifically:

$$\mu_t = \begin{cases} \max(\beta, l_{t_{1:M'}, d_{1:M}}), & \text{if } e_t^{\tau, \tau'} \leq 0 \\ \alpha, & \text{otherwise} \end{cases} \quad (9)$$

Since $l_{t_{1:M'}, d_{1:M}} \in [0, 1]$, the amount of optimism added by hysteretic updates is reduced; update will be taken in magnitudes proportional to the estimated likelihood.

5. Experiments

5.1. Recurrent architectures evaluation on meeting-in-a-grid

We conduct experiments using a Double-DQN architecture [14] as basis with added LSTM layer. The network starts with 2 fully connected layers of 32 and 64 neurons respectively, then a LSTM layer with 64 memory cells, a fully connected layer with 32 neurons which outputs value estimates for each action. We use $\beta = 0.4$, $\gamma = 0.95$ and Adam Optimizer [15] for training. For quantile estimators, we sample 16 τ and τ' to approximate return distribution for both training and acting online, and quantile embeddings were combined with the output of LSTM layer. Training is done in parallel, and results shown are batches of 20 randomly seeded runs. We use recurrence in meeting-in-a-grid domain to keep consistency with previous works [1]. The domain consists of one moving target and two agents in a grid world, agents get reward 1 for simultaneously standing on target location, 0 otherwise. Episode terminates after 40 transitions or upon successful meeting. Observation include probabilistically obscured locations of agent themselves and target, moving actions result in stochastic transitions.

We first evaluate RHIRQN performance against HDRQN and HIRQN. As shown in Figure 1, noticing differences in variance, that using quantile network alone or hysteretic-DQN alone not only fail to convergence to optimal policy on said benchmark, but also susceptible to divergence, producing lesser effective joint policy over time. [TODO: MAKE FIGURE AND TALK ABOUT LIKELIHOOD TREND DURING TRAINING, SHOW LIKELIHOOD DOMINATES HYSTERETIC OVERTIME AND ROBUST TO DIFFERENT HYSTERETIC VALUES]

5.2. Multi-Agent Object Transportation Problems (CMOTPs)

We evaluate our agents on three variations of CMOTPs [2, 16], consistent with Palmer et al.'s work on LDQN. Also grid-world tasks, CMOTPs requires two agents carrying a box to a desired location where agents get a terminal reward; the box only moves when agents are by its side and moving in the same direction. Different variations include more obstacles in the grid-world and stochastic terminal rewards. CMOTPs have 16 by 16 sensory observations with added



Figure 1. meeting-in-a-grid benchmark

noise.

Our network architecture is mostly same as LDQN for comparability: two convolutional layers with 32 and 64 kernels, a fully connected layers with 1024 neurons which combines quantile embedding, followed by another fully connected layer with 1024 neurons which then maps onto value estimates for each action. Noticing from Figure 2 that although both methods converge to joint optimal policy while our method show an improved sample efficiency. The temperature and leniency control parameters were brought from Palmer et al.’s work where they found most suitable for the task; we hypothesize that the temperature is decaying less aggressive than it should be, which is likely due to temperature folding techniques or the hashing space of the autoencoder is larger than needed.

On the other hand, our method which utilize time difference likelihood to guide negative updates, started to learn aggressively early on. Initially we worry that the Q estimates would be not optimistic enough to perform coordinated actions, but the likelihood estimates were able to produce small values in under-explored state-action space while not hesitate to update negatively in explored spaces.

References

- Claus, C. and Boutilier, C. The dynamics of reinforcement learning in cooperative multiagent systems. *AAAI/IAAI*, 1998:746–752, 1998.
- Fulda, N. and Ventura, D. Predicting and preventing coordination problems in cooperative q-learning systems. In *IJCAI*, volume 2007, pp. 780–785, 2007.
- Hausknecht, M. and Stone, P. Deep recurrent q-learning for partially observable mdps. *CoRR, abs/1507.06527*, 7(1), 2015.




Figure 2. CMOTP benchmark, aggregated three CMOTP variances

- Hochreiter, S. and Schmidhuber, J. Long short-term memory. *Neural computation*, 9(8):1735–1780, 1997.
- Lin, L.-J. Reinforcement learning for robots using neural networks. Technical report, Carnegie-Mellon Univ Pittsburgh PA School of Computer Science, 1993.
- Matignon, L., Laurent, G., and Le Fort-Piat, N. Hysteretic q-learning: an algorithm for decentralized reinforcement learning in cooperative multi-agent teams. In *IEEE/RSJ International Conference on Intelligent Robots and Systems, IROS’07.*, pp. 64–69, 2007.
- Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., Graves, A., Riedmiller, M., Fidjeland, A. K., Ostrovski, G., et al. Human-level control through deep reinforcement learning. *Nature*, 518(7540): 529, 2015.
- Omidshafiei, S., Pazis, J., Amato, C., How, J. P., and Vian, J. Deep decentralized multi-task multi-agent reinforcement learning under partial observability. *arXiv preprint arXiv:1703.06182*, 2017.
- Palmer, G., Tuyls, K., Bloembergen, D., and Savani, R. Lenient multi-agent deep reinforcement learning. In *Proceedings of the 17th International Conference on Autonomous Agents and MultiAgent Systems*, pp. 443–451. International Foundation for Autonomous Agents and Multiagent Systems, 2018.
- Panait, L., Sullivan, K., and Luke, S. Lenient learners in cooperative multiagent systems. In *Proceedings of the fifth international joint conference on Autonomous agents and multiagent systems*, pp. 801–803. ACM, 2006.
- Watkins, C. J. and Dayan, P. Q-learning. *Machine learning*, 8(3-4):279–292, 1992.